

hw3nk

September 19, 2024

```
[1]: import matplotlib.pyplot as plt
import pandas as pd
import numpy as np
import statsmodels.api as sm
import statsmodels.formula.api as smf
```

```
[3]: #Question 5 - 6.27
X = np.array([
    [1, 7, 33],
    [1, 4, 41],
    [1, 16, 7],
    [1, 3, 49],
    [1, 21, 5],
    [1, 8, 31]
])

y = np.array([42, 33, 75, 28, 91, 55])

#a
XTX_inv = np.linalg.inv(X.T @ X)
b = XTX_inv @ X.T @ y

print("(a) b:")
print(b)

#b
y_hat = X @ b
e = y - y_hat

print("\n(b) e:")
print(e)

#c
H = X @ XTX_inv @ X.T

print("\n(c) H:")
print(H)
```

```

#d
SSR = e.T @ e

print("\n(d) SSR:")
print(SSR)

#e
n, p = X.shape
s2 = SSR / (n - p + 1)
cov_b = s2 * XTX_inv

print("\n(e) Estimate of sigma^2:")
print(s2)

print("\n(e) cov_b:")
print(cov_b)

#f
X_h = np.array([1, 10, 30])
Y_h_hat = X_h @ b

print("\n(f) Y_h_hat:")
print(Y_h_hat)

#g
s2_Y_h = s2 * (X_h @ XTX_inv @ X_h.T)

print("\n(g) (s^2(Y_h):")
print(s2_Y_h)

```

(a) b:

```
[33.93210327  2.7847614 -0.26441893]
```

(b) e:

```
[-2.69960842 -1.22997279 -1.63735316 -1.32985996 -0.08999801  6.98679233]
```

(c) H:

```
[[ 0.23143293  0.25167585  0.21178735  0.14886839 -0.05475543  0.21099091]
 [ 0.25167585  0.31240459  0.09437844  0.26627729 -0.14787283  0.22313666]
 [ 0.21178735  0.09437844  0.70442026 -0.31917435  0.10446672  0.20412159]
 [ 0.14886839  0.26627729 -0.31917435  0.61425632  0.14143492  0.14833743]
 [-0.05475543 -0.14787283  0.10446672  0.14143492  0.94039955  0.01632707]
 [ 0.21099091  0.22313666  0.20412159  0.14833743  0.01632707  0.19708635]]
```

(d) SSR:

```
62.073538196057626
```

(e) Estimate of σ^2 :

15.518384549014407

(e) cov_b:

```
[[ 5.36603352e+02 -2.56191874e+01 -1.01962028e+01]
 [-2.56191874e+01  1.24624977e+00  4.83050531e-01]
 [-1.01962028e+01  4.83050531e-01  1.96850816e-01]]
```

(f) Y_h_hat:

53.847149399348694

(g) (s²(Y_h):

4.068464775116091

```
[6]: #Question 1
insurancedata=pd.read_csv('../data/insurance.csv')
insurancedata.sample(10)
```

```
[6]:      age    sex    bmi  children  smoker    region    charges
1328   23  female  24.225         2     no  northeast  22395.74424
670    30   male  31.570         3     no  southeast   4837.58230
524    42   male  26.070         1    yes  southeast  38245.59327
1133   52  female  18.335         0     no  northwest   9991.03765
741    27   male  29.150         0    yes  southeast  18246.49550
606    27  female  25.175         0     no  northeast   3558.62025
347    46   male  33.345         1     no  northeast   8334.45755
624    59   male  28.785         0     no  northwest  12129.61415
92     59   male  29.830         3    yes  northeast  30184.93670
1334   18  female  31.920         0     no  northeast   2205.98080
```

```
[8]: insurance=insurancedata[['charges','age','bmi','children']].copy()
insurance.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1338 entries, 0 to 1337
Data columns (total 4 columns):
#   Column      Non-Null Count  Dtype
---  -
0   charges     1338 non-null   float64
1   age         1338 non-null   int64
2   bmi         1338 non-null   float64
3   children    1338 non-null   int64
dtypes: float64(2), int64(2)
memory usage: 41.9 KB
```

```
[11]: reg = smf.ols('charges~age+bmi+children',data=insurance).fit()

# print(dir(reg)) #members of the object provided by the modelling
```

```
[12]: reg.summary()
```

```
[12]:
```

Dep. Variable:	charges	R-squared:	0.120
Model:	OLS	Adj. R-squared:	0.118
Method:	Least Squares	F-statistic:	60.69
Date:	Thu, 19 Sep 2024	Prob (F-statistic):	8.80e-37
Time:	16:47:18	Log-Likelihood:	-14392.
No. Observations:	1338	AIC:	2.879e+04
Df Residuals:	1334	BIC:	2.881e+04
Df Model:	3		
Covariance Type:	nonrobust		

	coef	std err	t	P> t	[0.025	0.975]
Intercept	-6916.2433	1757.480	-3.935	0.000	-1.04e+04	-3468.518
age	239.9945	22.289	10.767	0.000	196.269	283.720
bmi	332.0834	51.310	6.472	0.000	231.425	432.741
children	542.8647	258.241	2.102	0.036	36.261	1049.468

Omnibus:	325.395	Durbin-Watson:	2.012
Prob(Omnibus):	0.000	Jarque-Bera (JB):	603.372
Skew:	1.520	Prob(JB):	9.54e-132
Kurtosis:	4.255	Cond. No.	290.

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

```
[13]: sm.stats.anova_lm(reg,typ=1)
```

```
[13]:
```

	df	sum_sq	mean_sq	F	PR(>F)
age	1.0	1.753019e+10	1.753019e+10	135.546341	6.627851e-30
bmi	1.0	5.446449e+09	5.446449e+09	42.112843	1.211545e-10
children	1.0	5.715190e+08	5.715190e+08	4.419080	3.572625e-02
Residual	1334.0	1.725261e+11	1.293299e+08	NaN	NaN

```
[14]: sm.stats.anova_lm(reg,typ=2)
```

```
[14]:
```

	sum_sq	df	F	PR(>F)
age	1.499426e+10	1.0	115.938067	5.533923e-26
bmi	5.417280e+09	1.0	41.887301	1.354882e-10
children	5.715190e+08	1.0	4.419080	3.572625e-02
Residual	1.725261e+11	1334.0	NaN	NaN

```
[15]: #Question 2
```

```
column_names = ['Y', 'X1', 'X2', 'X3', 'X4']
propertydata = pd.read_csv('../data/property.txt', delim_whitespace=True,
                             names=column_names)

propertydata.info()
```

```
<class 'pandas.core.frame.DataFrame'>
```

RangeIndex: 81 entries, 0 to 80

Data columns (total 5 columns):

#	Column	Non-Null Count	Dtype
0	Y	81 non-null	float64
1	X1	81 non-null	int64
2	X2	81 non-null	float64
3	X3	81 non-null	float64
4	X4	81 non-null	int64

dtypes: float64(3), int64(2)

memory usage: 3.3 KB

/var/folders/86/c2gz31wn29b2r_53d_q3g3hc0000gn/T/ipykernel_66472/396240547.py:2:

FutureWarning: The 'delim_whitespace' keyword in pd.read_csv is deprecated and will be removed in a future version. Use ``sep='\s+'`` instead

```
propertydata = pd.read_csv('../data/property.txt', delim_whitespace=True,
names=column_names)
```

```
[32]: #a
print("(a)")
model = smf.ols('Y ~ X1 + X2 + X3 + X4', data=propertydata).fit()
print(model.summary())
```

(a)

OLS Regression Results

```
=====
Dep. Variable:          Y      R-squared:                0.585
Model:                  OLS    Adj. R-squared:           0.563
Method:                 Least Squares    F-statistic:      26.76
Date:                   Thu, 19 Sep 2024    Prob (F-statistic):  7.27e-14
Time:                   20:27:27    Log-Likelihood:     -122.75
No. Observations:      81    AIC:                    255.5
Df Residuals:          76    BIC:                    267.5
Df Model:               4
Covariance Type:       nonrobust
=====
```

	coef	std err	t	P> t	[0.025	0.975]
Intercept	12.2006	0.578	21.110	0.000	11.049	13.352
X1	-0.1420	0.021	-6.655	0.000	-0.185	-0.100
X2	0.2820	0.063	4.464	0.000	0.156	0.408
X3	0.6193	1.087	0.570	0.570	-1.545	2.784
X4	7.924e-06	1.38e-06	5.722	0.000	5.17e-06	1.07e-05

```
=====
Omnibus:                1.922    Durbin-Watson:           1.580
Prob(Omnibus):           0.383    Jarque-Bera (JB):         1.301
Skew:                    0.148    Prob(JB):                 0.522
Kurtosis:                3.545    Cond. No.                 1.74e+06
=====
```

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

[2] The condition number is large, 1.74e+06. This might indicate that there are strong multicollinearity or other numerical problems.

```
[34]: #b
print("(b)")
fitted_values = model.fittedvalues
residuals = model.resid
print("\nFitted Values:")
print(fitted_values.head(6))
print("\nResiduals:")
print(residuals.head(6))
sigma_squared = (residuals ** 2).sum() / (len(propertydata) - len(model.params))
print("\nSigma^2:")
print(sigma_squared)
```

(b)

Fitted Values:

```
0    14.535672
1    13.513806
2    11.091053
3    15.133568
4    13.686716
5    13.687185
dtype: float64
```

Residuals:

```
0    -1.035672
1    -1.513806
2    -0.591053
3    -0.133568
4     0.313284
5    -3.187185
dtype: float64
```

Sigma^2:

```
1.2925078150380076
```

```
[35]: #c
print("(c)")
p_val_b2 = model.pvalues['X2']
beta_2 = model.params['X2']
print("\nPval B2:")
print(p_val_b2)
```

```
print("\nEstimate B2:")
print(beta_2)
```

(c)

Pval B2:
2.747396037799102e-05

Estimate B2:
0.28201652995092874

```
[41]: #d
print("(d)")
type1 = sm.stats.anova_lm(model, typ=1)
print("\nType 1 results:")
print(type1)

p_val_t1_b1 = type1.loc['X1', 'PR(>F)']
p_val_t1_b2 = type1.loc['X2', 'PR(>F)']

print("\nPval Beta 1:")
print(p_val_t1_b1)
print("\nPval Beta 2:")
print(p_val_t1_b2)

type2 = sm.stats.anova_lm(model, typ=2)
print("\nType 2 results:")
print(type2)

p_val_t2_b1 = type2.loc['X1', 'PR(>F)']
p_val_t2_b2 = type2.loc['X2', 'PR(>F)']

print("\nPval Beta 1:")
print(p_val_t2_b1)
print("\nPval Beta 2:")
print(p_val_t2_b2)
```

(d)

Type 1 results:

	df	sum_sq	mean_sq	F	PR(>F)
X1	1.0	14.818520	14.818520	11.464936	1.125291e-03
X2	1.0	72.802011	72.802011	56.326167	9.699085e-11
X3	1.0	8.381417	8.381417	6.484616	1.290389e-02
X4	1.0	42.324958	42.324958	32.746385	1.975990e-07
Residual	76.0	98.230594	1.292508	NaN	NaN

Pval Beta 1:

0.0011252906812985895

Pval Beta 2:

9.6990847500436e-11

Type 2 results:

	sum_sq	df	F	PR(>F)
X1	57.242762	1.0	44.288136	3.894322e-09
X2	25.758955	1.0	19.929439	2.747396e-05
X3	0.419746	1.0	0.324753	5.704457e-01
X4	42.324958	1.0	32.746385	1.975990e-07
Residual	98.230594	76.0	NaN	NaN

Pval Beta 1:

3.894321773978801e-09

Pval Beta 2:

2.7473960377990986e-05

```
[39]: #e
print("(e)")
new_property = pd.DataFrame({'X1': [4], 'X2': [10], 'X3': [0.1], 'X4': [80000]})

predictions = model.get_prediction(new_property)
prediction_summary = predictions.summary_frame(alpha=0.10)

print("\nPrediction Interval 90%")
print(prediction_summary[['obs_ci_lower', 'obs_ci_upper']])
```

(e)

Prediction Interval 90%

	obs_ci_lower	obs_ci_upper
0	13.228907	17.068083